

Logistic Regression

Guowei Wei
Department of Mathematics
Michigan State University

References:


Duc D. Nguyen's lecture notes

Andrew Ng's notes

Wikipedia

Introduction

Statistical classification: identifying to which of a set of categories a new observation belongs, on the basis of a training set of data.



Introduction

- Classification
- Examples:
 - Matter: Toxic / Not Toxic ?
 - Students: Pass / Fail ?
 - Van Gogh Painting: Real / Fake?
- Labels $y \in \{0,1\}$. 0: Negative class (Fail), 1: Positive class (Pass)
- We can have more than two labels $y \in \{0,1,2,3\}$ for DNA (cytosine [C], guanine [G], adenine [A] and thymine [T])


Example

- Performance of group of 20 students spend between 0 and 6 hours studying for an exam

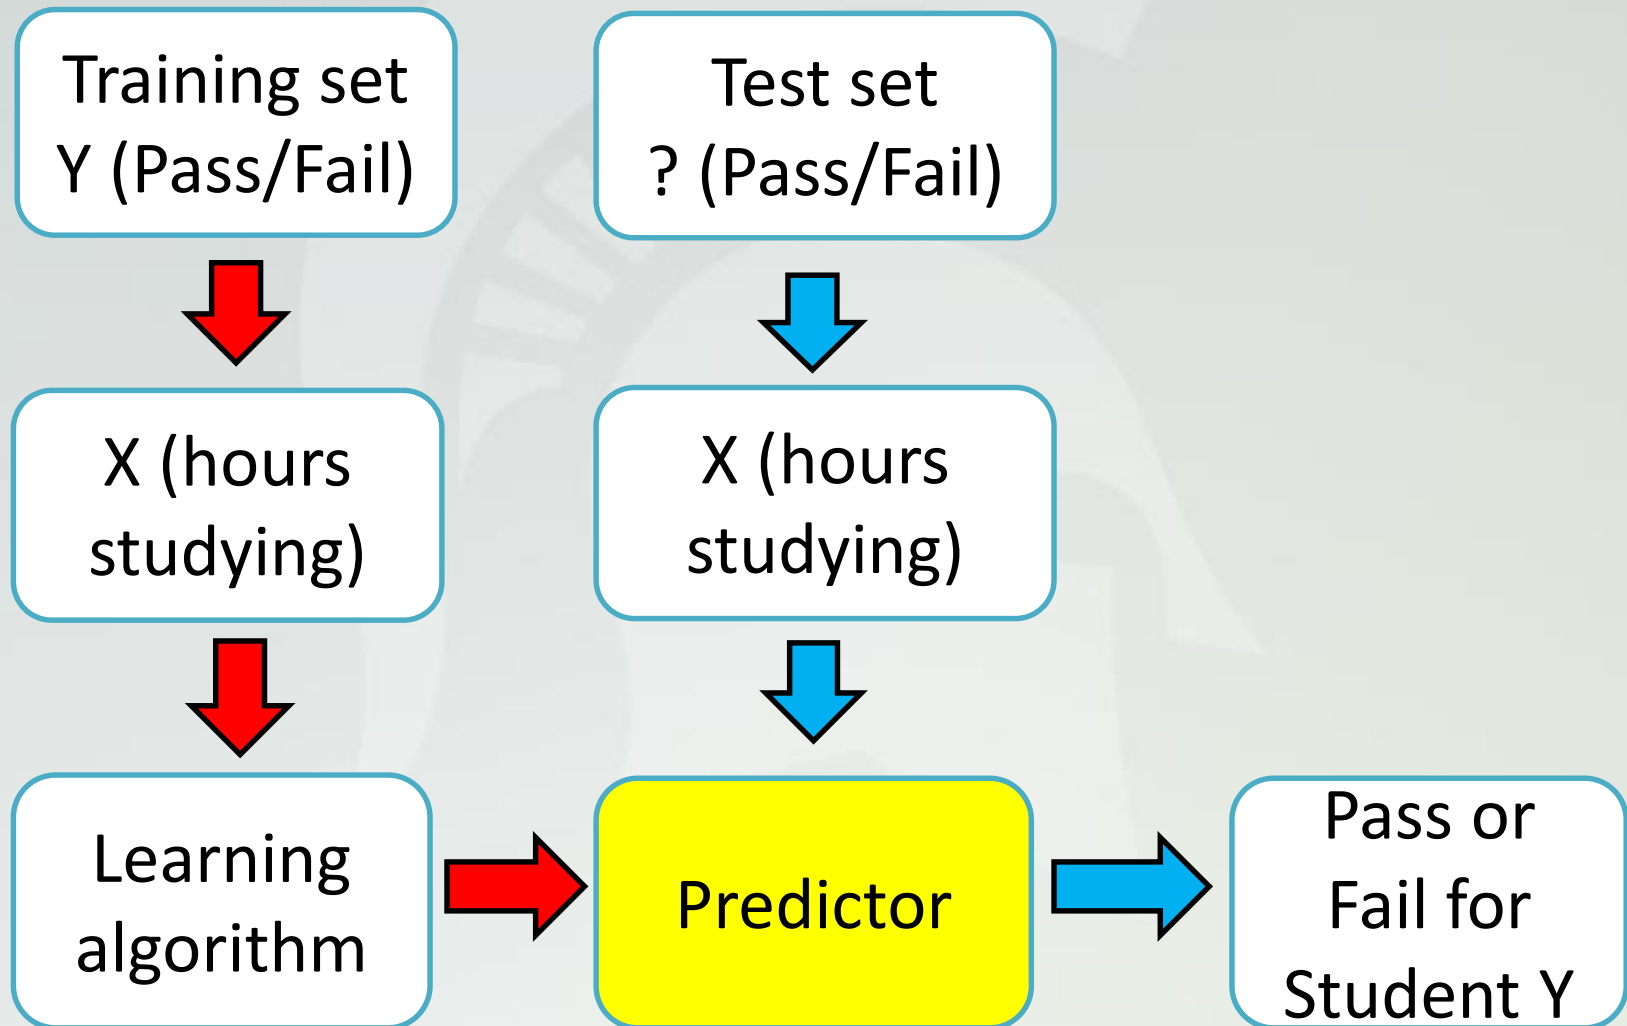
Hours	Pass/Fail
0.5	Fail
0.75	Fail
1.75	Pass
2.25	Pass
...	...

Example

Original data



Model Representation



Predictor Construction

- Wish to construct a predictor:

$$p_c(\mathbf{x}) = ?$$


- Our labels only has two values 0 and 1
- Predictor always give a real value.
- $0 \leq p_c(\mathbf{x}) \leq 1$
- To get classification, choose a threshold z :

$$p_c(\mathbf{x}) < z \text{ then } y = 0$$

$$p_c(\mathbf{x}) \geq z \text{ then } y = 1$$

- With two labels $\{0,1\}$, it is natural to choose $z = 0.5$

Predictor Construction



Predictor Construction

- We choose

Sigmoid / Logistic function:


$$p_{\mathbf{c}}(\mathbf{x}) = \frac{1}{1 + e^{-\mathbf{c}^T \mathbf{x}}}$$



(Is this model still linear?)

- If dataset has one feature, i.e., $\mathbf{x} = (1, x_1)^T$, then $\mathbf{c} = (c_0, c_1)^T$

Predictor Construction



Predictor Construction

- If dataset has n features, i.e., $\mathbf{x} = (1, x_1, \dots, x_n)^T$, then $\mathbf{c} = (c_0, c_1, \dots, c_n)^T$
- Single-layer perceptron

$$p_{\mathbf{c}}(\mathbf{x}) = \frac{1}{1 + e^{-c_0 - c_1 x_1 - c_2 x_2 - \dots - c_n x_n}}$$

or

$$p_{\mathbf{c}}(\mathbf{x}) = \frac{1}{1 + e^{-\mathbf{c}^T \mathbf{x}}}$$

is the probability that $y = 1$ on input \mathbf{x}

Predictor Construction

- We can also use the polynomial logistic function, e.g.,

$$p_{\mathbf{c}}(\mathbf{x}) = \frac{1}{1 + e^{-c_0 - c_1 x_1 - c_2 x_1^2}} \quad \text{or more general:}$$

$$p_{\mathbf{c}}(\mathbf{x}) = \frac{1}{1 + e^{-c_0 - c_{11}x_1 - c_{12}x_1^2 + \dots, c_{n1}x_n - c_{n2}x_n^2 + \dots}}$$

- **Loss function construction:**

- Linear regression:

$$L(\mathbf{c}) = \sum_{i=1}^m (p(x^{(i)}) - y^{(i)})^2$$

Predictor Construction

- We can also use the polynomial logistic function, e.g.,

$$p_{\mathbf{c}}(\mathbf{x}) = \frac{1}{1 + e^{-c_0 - c_1 x_1 - c_2 x_1^2}} \quad \text{or more general:}$$

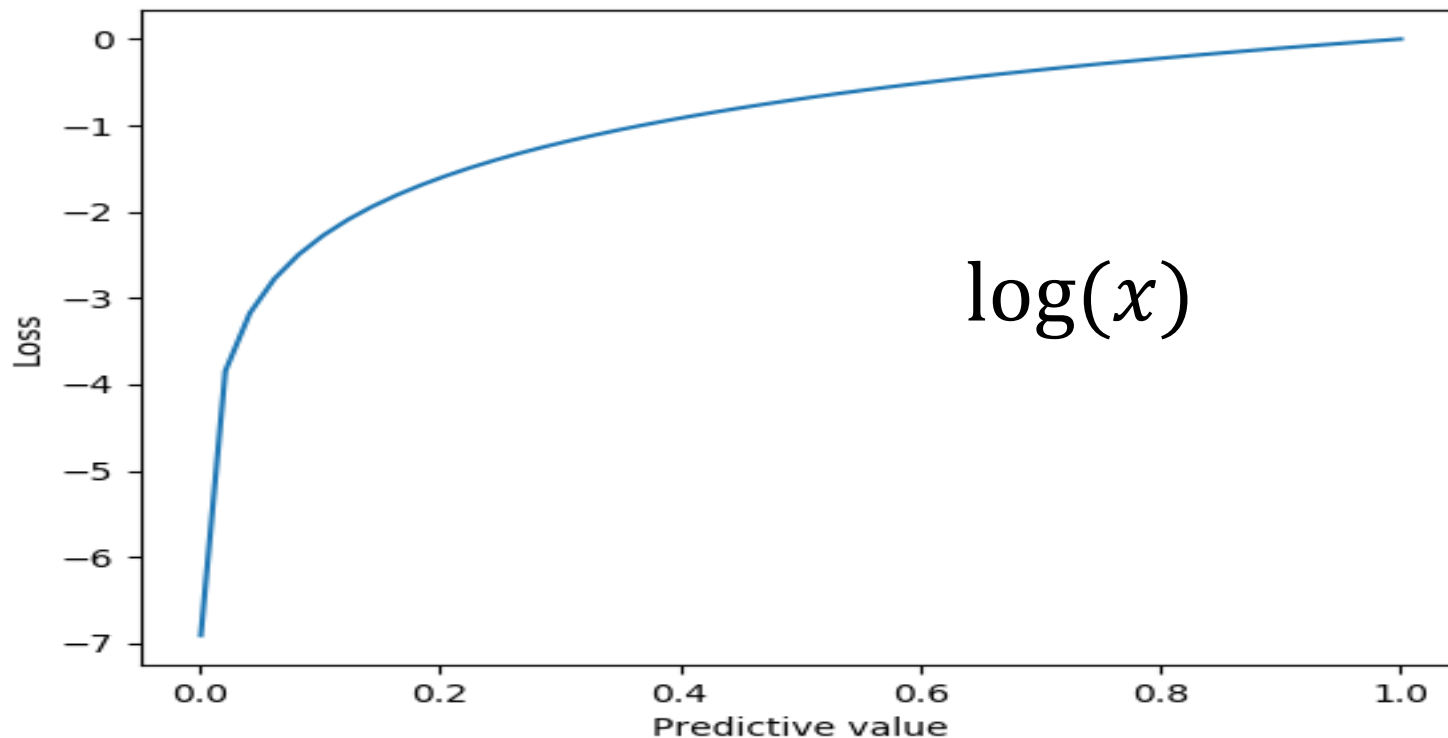
$$p_{\mathbf{c}}(\mathbf{x}) = \frac{1}{1 + e^{-c_0 - c_{11}x_1 - c_{12}x_1^2 + \dots, c_{n1}x_n - c_{n2}x_n^2 + \dots}}$$

$$p_{\mathbf{c}}(\mathbf{x}) = \frac{1}{1 + e^{-f(\mathbf{c}, \mathbf{x})}}$$

How to determine coefficients?


Loss function construction

- When $y = 1$, higher $p_c(\mathbf{x})$ is more accurate \Rightarrow smaller lost
- Any function reflects that behavior?



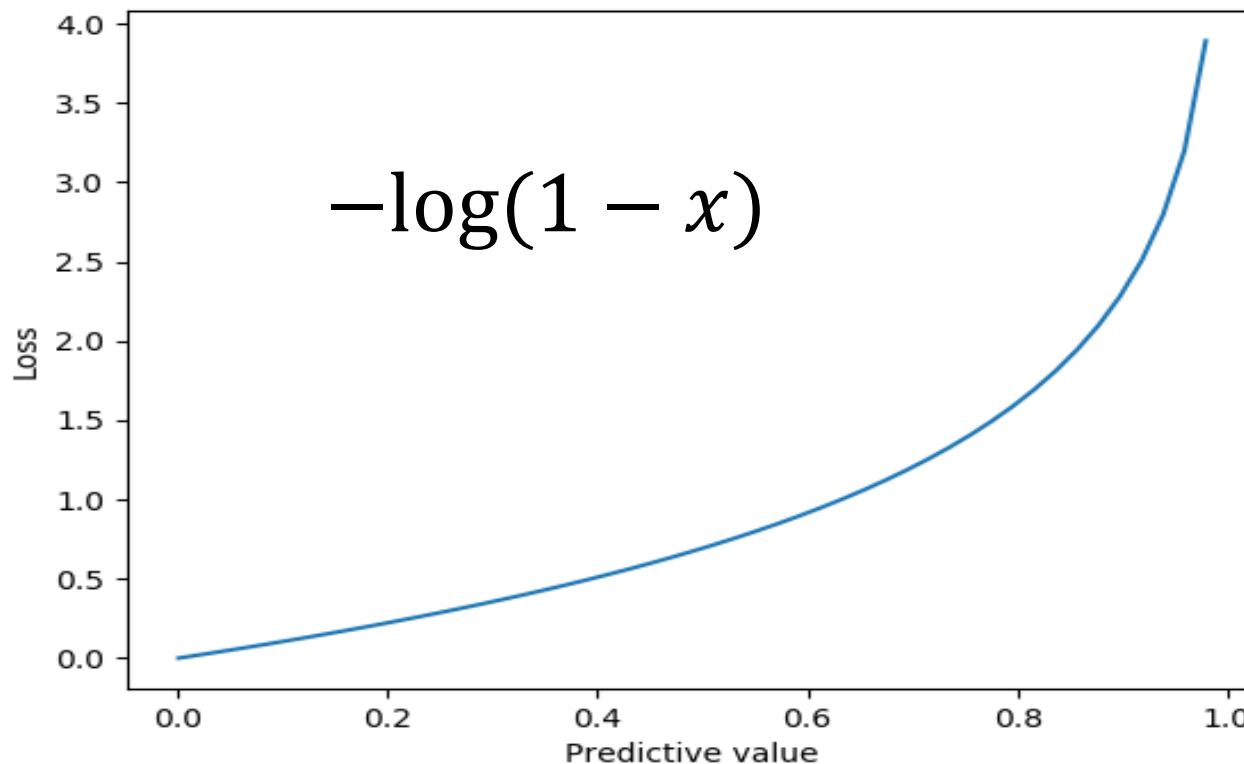
Loss function construction

Make it nicer



Loss function construction

- When $y = 0$, lower $p_c(\mathbf{x})$ is more accurate \Rightarrow smaller loss
- Any function reflects that behavior?



Loss function construction

- So we have

$$L(p_c(\mathbf{x}), y) = \begin{cases} -\log(p_c(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - p_c(\mathbf{x})) & \text{if } y = 0 \end{cases}$$

- Combine two goals into one function:


$$\begin{aligned} L(p_c(\mathbf{x}), y) \\ = -y \log(p_c(\mathbf{x})) - (1 - y) \log(1 - p_c(\mathbf{x})) \end{aligned}$$

Loss function

- Total loss for the whole dataset


$$\begin{aligned} L(\mathbf{c}) \\ &= \frac{1}{m} \sum_{i=1}^m \left[-y^{(i)} \log \left(p_{\mathbf{c}}(\mathbf{x}^{(i)}) \right) \right] \end{aligned}$$

Resulting Predicted Model



Decision Boundary

- Boundary that separates this class from another ones
- We can construct the **decision boundary** from optimal predictor



$$p_c(\mathbf{x}) = \frac{1}{1 + e^{-\mathbf{c}^T \mathbf{x}}}$$

$$\mathbf{c}^T = (-1, 1, 1)$$

$$\mathbf{x}^T = (1, x_1, x_2)$$

$$p_c(\mathbf{x}) = \frac{1}{1 + e^{1-x_1-x_2}}$$

When $p_c(\mathbf{x}) = 0.5$ (threshold)

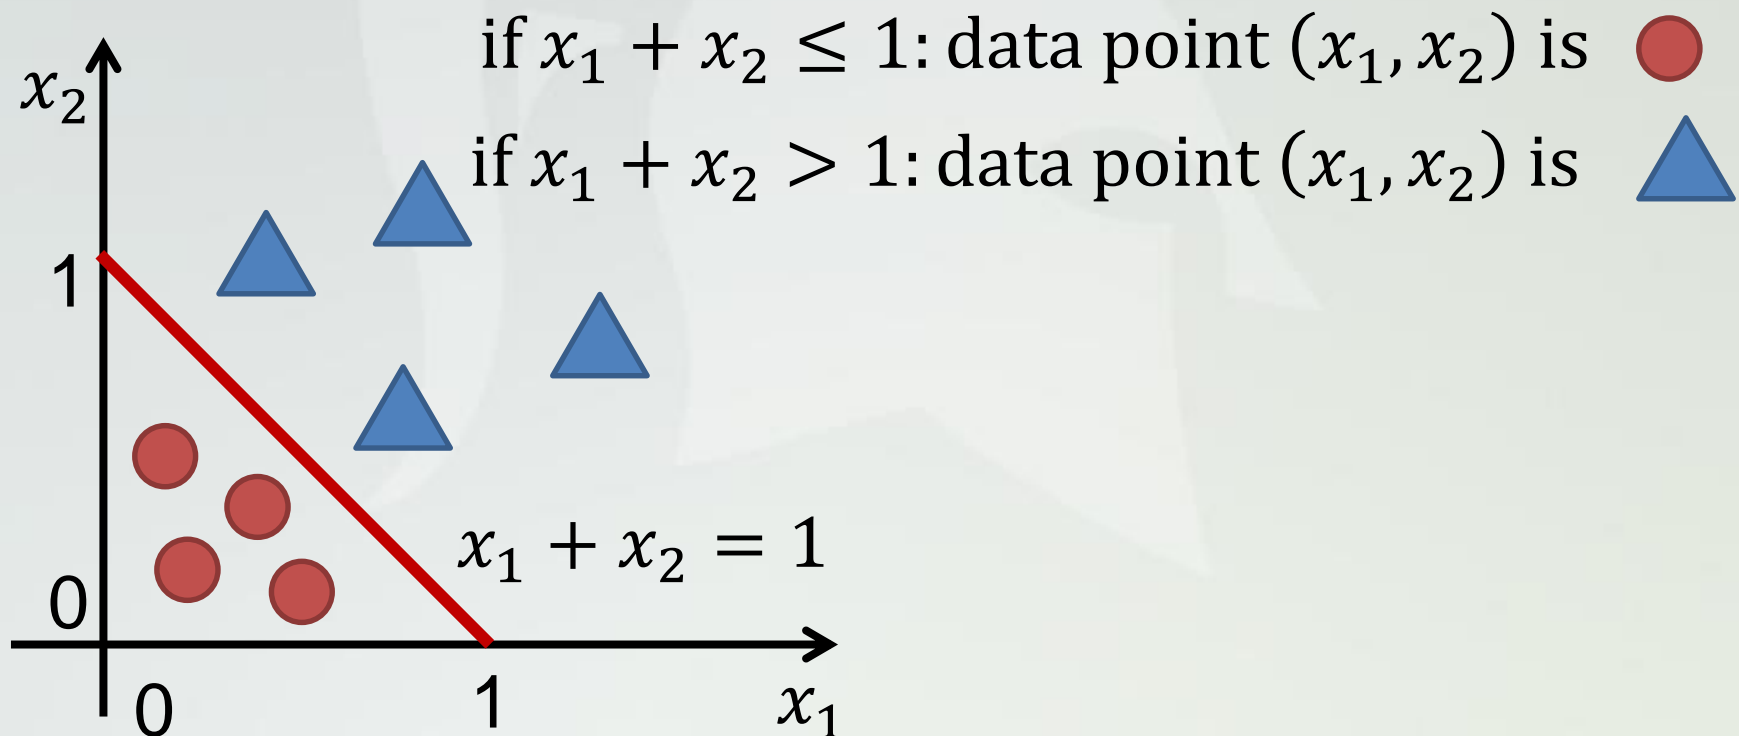
$$\frac{1}{1 + e^{1-x_1-x_2}} = \frac{1}{2}$$

$$\Rightarrow x_1 + x_2 = 1$$

Decision boundary


Decision Boundary

- Boundary that separates this class from another ones
- We can construct the decision boundary from optimal predictor



Decision Boundary

- Decision boundary can be a curve or manifold



$$p_c(\mathbf{x}) = \frac{1}{1 + e^{1-x_1^2-x_2^2}}$$
$$p_c(\mathbf{x}) = 0.5$$
$$\frac{1}{1 + e^{1-x_1^2-x_2^2}} = \frac{1}{2}$$
$$\Rightarrow x_1^2 + x_2^2 = 1$$

Multi-Class

- Loss function designs for two classes
- What if we have more than 2 classes?
- Examples:
 - Student Performance: A, B, C, D, F
 - Weather prediction: Sunny, Cloudy, Rain, Snow

Multi-Class

- One-vs-all:
 - If we have n classes: l_1, l_2, \dots, l_n
 - For each class l_i . Consider two labels:
 l_i and not l_i
 - Train logistic regression classifier for this case to get the probability $p_{\mathbf{c}}^{l_i}(\mathbf{x})$ that $y = l_i$
 - Pick the class l_k that maximizes
 $\max_k p_{\mathbf{c}}^{l_k}(\mathbf{x})$